

Note to other teachers and users of these slides: We would be delighted if you found our material useful for giving your own lectures. Feel free to use these slides verbatim, or to modify them to fit your own needs. If you make use of a significant portion of these slides in your own lecture, please include this message, or a link to our web site: <http://www.mmds.org>

Learning through Experimentation

CS246: Mining Massive Datasets
Jure Leskovec, Stanford University
<http://cs246.stanford.edu>



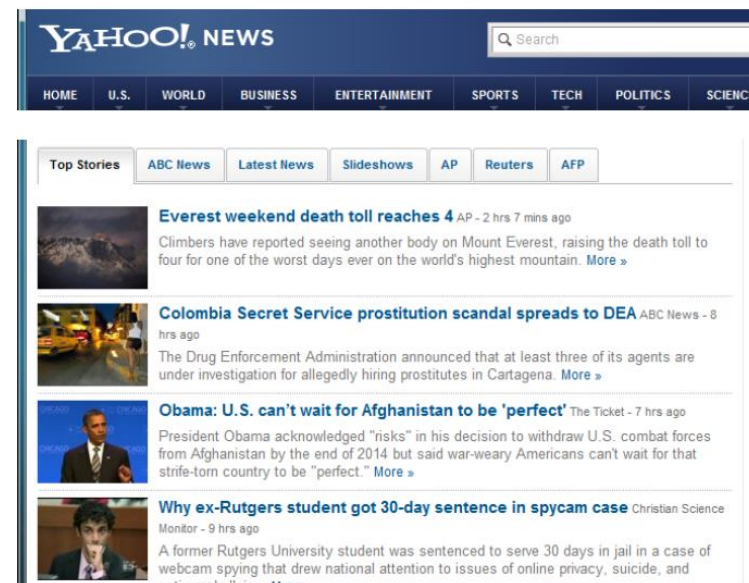
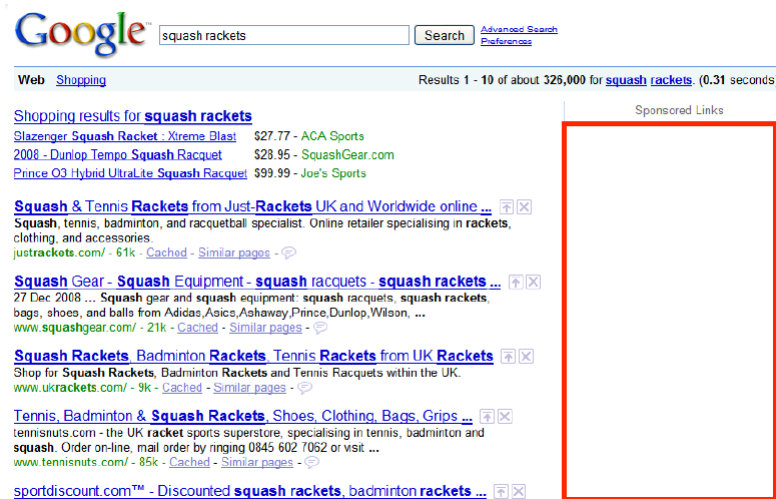
Learning through Experimentation

■ Web advertising

- We discussed how to match advertisers to queries in real-time
- But we did not discuss how to estimate the **CTR** (Click-Through Rate)

■ Recommendation engines

- We discussed how to build recommender systems
- But we did not discuss the **cold-start** problem



Learning through Experimentation

- What do **CTR** and **cold-start** have in common?
- With every **ad we show/product we recommend** we gather more data about the **ad/product**
- Theme: Learning through experimentation

A screenshot of a Google search results page for the query "squash rackets". The page shows several organic search results from various websites like "Just-Rackets UK", "SquashGear.com", and "TennisRackets.com". On the right side of the page, there is a red rectangular box highlighting the "Sponsored Links" area, which is currently empty.

A screenshot of a Yahoo! News page. The page features a navigation bar with categories like HOME, U.S., WORLD, BUSINESS, ENTERTAINMENT, SPORTS, TECH, POLITICS, and SCIENCE. Below the navigation bar, there are tabs for "Top Stories", "ABC News", "Latest News", "Slideshows", "AP", "Reuters", and "AFP". The main content area displays several news stories with accompanying images and headlines, such as "Everest weekend death toll reaches 4", "Colombia Secret Service prostitution scandal spreads to DEA", "Obama: U.S. can't wait for Afghanistan to be 'perfect'", and "Why ex-Rutgers student got 30-day sentence in spycam case".

Example: Web Advertising

- **Google's goal: Maximize revenue**
- **The old way: Pay by impression (CPM)**
 - **Best strategy: Go with the highest bidder**
 - But this ignores the “effectiveness” of an ad
- **The new way: Pay per click! (CPC)**
 - **Best strategy: Go with expected revenue**
 - What's the expected revenue of ad a for query q ?
 - $E[\text{revenue}_{a,q}] = P(\text{click}_a \mid q) * \text{amount}_{a,q}$

Prob. user will click on ad a given
that she issues query q
(Unknown! Need to gather information)

Bid amount for
ad a on query q
(Known)

Other Applications

- **Clinical trials:**

- Investigate effects of different treatments while minimizing adverse effects on patients

- **Adaptive routing:**

- Minimize delay in the network by investigating different routes

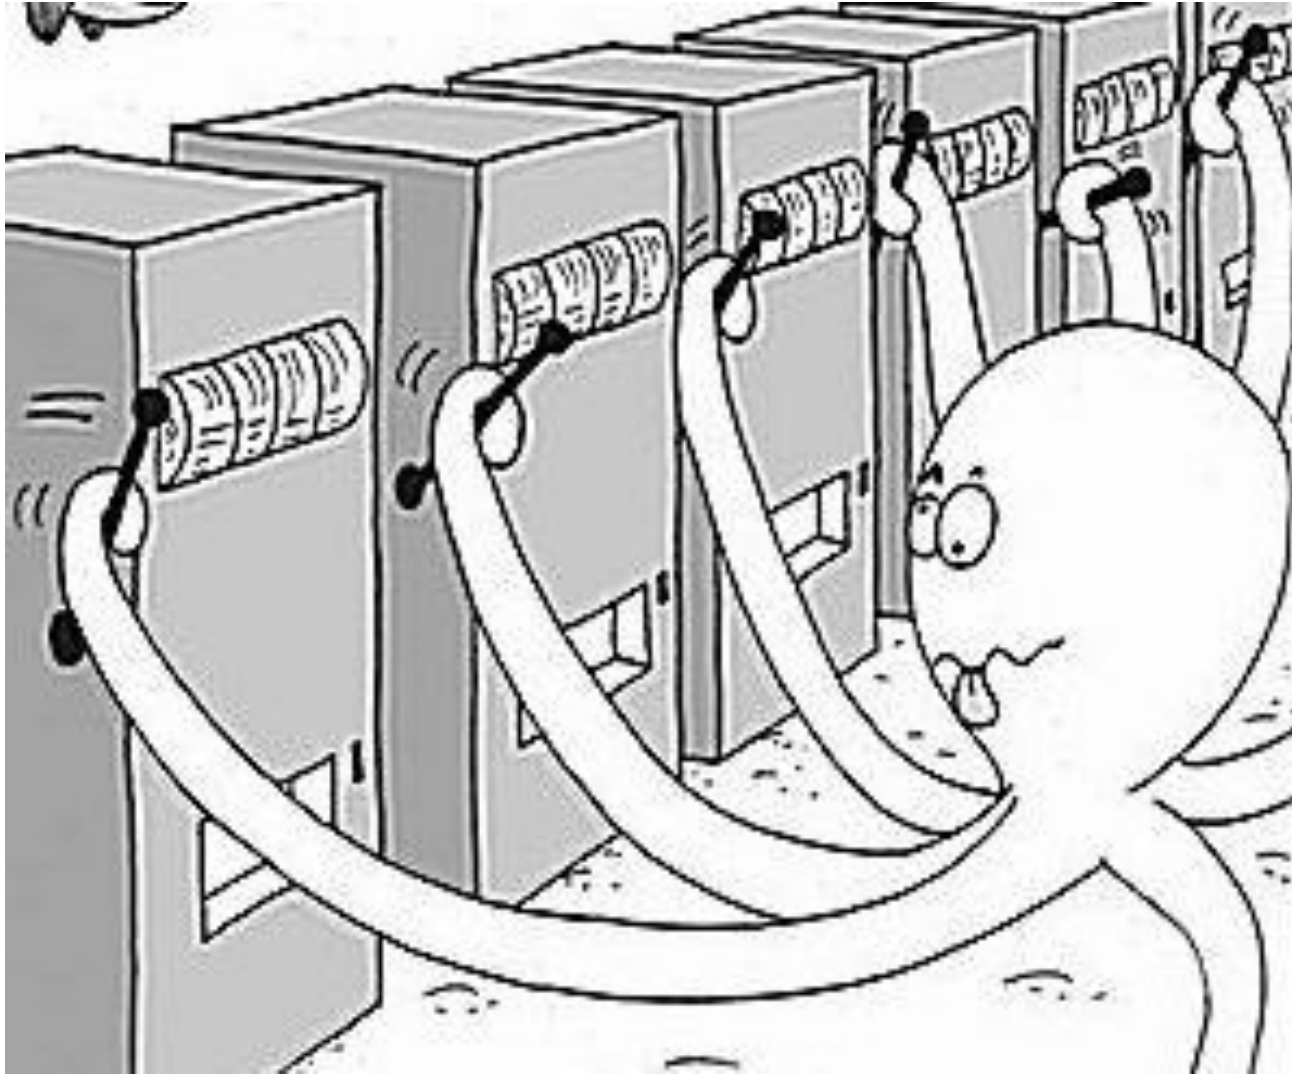
- **Asset pricing:**

- Figure out product prices while trying to make most money


Approach: Bandits



Approach: Multiarmed Bandits




k-Armed Bandit



- Each arm a
 - Wins (reward=1) with fixed (unknown) prob. μ_a
 - Loses (reward=0) with fixed (unknown) prob. $1-\mu_a$
- All draws are independent given $\mu_1 \dots \mu_k$
- How to pull arms to maximize total reward?

k-Armed Bandit



- How does this map to our setting?
- Each **query** is a **bandit**
- Each **ad** is an **arm**
- We want to estimate μ_a , the arm's probability of winning (i.e., ad's CTR μ_a)
- Every time we pull an arm we do an 'experiment'

Stochastic k-Armed Bandit

The setting:

- Set of k choices (arms)
- Each choice a is associated with unknown probability distribution P_a supported in $[0,1]$
- We play the game for T rounds
- In each round t :
 - (1) We pick some arm a
 - (2) We obtain random sample X_t from P_a
 - Note reward is independent of previous draws
- Our goal is to maximize $\sum_{t=1}^T X_t$
- Problem: we don't know μ_a ! But every time we pull some arm a we get to learn a bit about μ_a

Online Optimization

■ Online optimization with limited feedback

Choices	X_1	X_2	X_3	X_4	X_5	X_6	...
a_1					1	1	
a_2	0		1	0			
...							
a_k		0					

Time →

■ Like in online algorithms:

- Have to make a choice each time
- But we only receive information about the chosen action

Solving the Bandit Problem

- **Policy:** a strategy/rule that tells me which arm to pull in each iteration
 - Hopefully policy depends on the history of rewards
- **How to quantify performance of the algorithm? Regret!**

Performance Metric: Regret

- Let μ_a be the mean reward of P_a
- Payoff/reward of **best arm**: $\mu^* = \max_a \mu_a$
- Let $i_1, i_2 \dots i_T$ be the sequence of arms pulled
- **Instantaneous regret** at time t : $r_t = \mu^* - \mu_{i_t}$
- **Total regret**:

$$R_T = \sum_{t=1}^T r_t$$

- Typical goal: **Want a policy (arm allocation strategy) that guarantees: $\frac{R_T}{T} \rightarrow 0$ as $T \rightarrow \infty$**

- Note: Ensuring $R_T/T \rightarrow 0$ is stronger than maximizing payoffs (minimizing regret), as it means that in the limit we discover the true best hand.

Allocation Strategies

- If we knew the payoffs, which arm would we pull?

Pick $\arg \max_a \mu_a$

- What if we only care about estimating payoffs μ_a ?

- Pick each of k arms equally often: $\frac{T}{k}$

- **Estimate:** $\widehat{\mu}_a = \frac{k}{T} \sum_{j=1}^{T/k} X_{a,j}$

- **Regret:** $R_T = \frac{T}{k} \sum_{a=1}^k (\mu^* - \widehat{\mu}_a)$

$X_{a,j}$... payoff received
when pulling arm a for
 j -th time

Bandit Algorithm: First try

- Regret is defined in terms of average reward
- So, if we can estimate avg. reward we can minimize regret
- Consider algorithm: *Greedy*
Take the action with the highest avg. reward
 - **Example:** Consider 2 actions
 - **A1** reward 1 with prob. 0.3
 - **A2** has reward 1 with prob. 0.7
 - Play **A1**, get reward 1
 - Play **A2**, get reward 0
 - Now avg. reward of **A1** will never drop to 0, and we will never play action **A2**

Exploration vs. Exploitation

- The example illustrates a classic problem in **decision making**:
 - We need to trade off between **exploration** (gathering data about arm payoffs) and **exploitation** (making decisions based on data already gathered)
- The Greedy algo does not explore sufficiently
 - **Exploration**: Pull an arm we never pulled before
 - **Exploitation**: Pull an arm a for which we currently have the highest estimate of μ_a

Optimism

- The problem with our **Greedy** algorithm is that it is **too certain** in the estimate of μ_a
 - When we have seen a single reward of 0 we shouldn't conclude the average reward is 0
- **Greedy can converge to a suboptimal solution!**

New Algorithm: Epsilon-Greedy

Algorithm: Epsilon-Greedy

■ For $t=1:T$

- Set $\epsilon_t = O\left(\frac{1}{t}\right)$ (that is, ϵ_t decays over time t as $1/t$)
- **With prob. ϵ_t : Explore** by picking an arm chosen uniformly at random
- **With prob. $1 - \epsilon_t$: Exploit** by picking an arm with highest empirical mean payoff

■ Theorem [Auer et al. '02]

For suitable choice of ϵ_t it holds that

$$R_T = O(k \log T) \Rightarrow \frac{R_T}{T} = O\left(\frac{k \log T}{T}\right) \rightarrow 0$$

k ...number
of arms

Issues with Epsilon-Greedy

- What are some issues with **Epsilon-Greedy**?
 - **“Not elegant”**: Algorithm explicitly distinguishes between exploration and exploitation
 - **More importantly**: Exploration makes **suboptimal choices** (since it picks any arm equally likely)
- **Idea**: When exploring/exploiting we need to **compare** arms

Comparing Arms

- **Suppose we have done experiments:**
 - Arm 1: 1 0 0 1 1 0 0 1 0 1
 - Arm 2: 1
 - Arm 3: 1 1 0 1 1 1 0 1 1 1
- **Mean arm values:**
 - Arm 1: 5/10, Arm 2: 1, Arm 3: 8/10
- **Which arm would you pick next?**
- **Idea: Don't just look at the mean (that is, expected payoff) but also the confidence!**


Confidence Intervals (1)

- A confidence interval is a range of values within which we are sure the mean lies with a certain probability
 - We could believe μ_a is within $[0.2, 0.5]$ with probability 0.95
 - If we would have tried an action less often, our estimated reward is less accurate so the confidence interval is larger
 - Interval shrinks as we get more information (try the action more often)

Confidence Intervals (2)

- Assuming we know the confidence intervals
- Then, instead of **trying the action with the highest mean** we can **try the action with the highest upper bound on its confidence interval**
- This is called an **optimistic policy**
 - We believe an action is as good as possible given the available evidence

Confidence Based Selection



Calculating Confidence Bounds

Suppose we fix arm a :

- Let $X_{a,1} \dots X_{a,m}$ be the payoffs of arm a in the first m trials
 - So, $X_{a,1} \dots X_{a,m}$ are i.i.d. rnd. vars. taking values in $[0,1]$
- **Mean payoff of arm a :** $\mu_a = E[X_{a,\cdot}]$
- **Our estimate:** $\widehat{\mu}_{a,m} = \frac{1}{m} \sum_{\ell=1}^m X_{a,\ell}$
- Want to find b such that with high probability $|\mu_a - \widehat{\mu}_{a,m}| \leq b$
 - Want b to be as small as possible (so our estimate is close)
- **Goal:** Want to bound $P(|\mu_a - \widehat{\mu}_{a,m}| \leq b)$

Hoeffding's Inequality (1)

Hoeffding's inequality provides an upper bound on the probability that the average deviates from its expected value by more than a certain amount:

- Let $X_1 \dots X_m$ be i.i.d. rnd. vars. taking values in $[0,1]$
- Let $\mu = E[X]$ and $\hat{\mu}_m = \frac{1}{m} \sum_{\ell=1}^m X_\ell$
- **Then:** $P(|\mu - \hat{\mu}_m| \geq b) \leq 2 \exp(-2b^2m) = \delta$
 - δ ... is the confidence level
- **To find out the confidence interval b (for a given confidence level δ) we solve:**
 - $2e^{-2b^2m} \leq \delta$ then $-2b^2m \leq \ln(\delta/2)$
 - **So:** $b \geq \sqrt{\frac{\ln(\frac{2}{\delta})}{2m}}$

Hoeffding's Inequality (2)

- $P(|\mu - \hat{\mu}_m| \geq b) \leq 2 \exp(-2b^2 m)$
where b is our upper bound, m is number of times we played the action
- Let's set $b = b(a, T) = \sqrt{2 \log(T) / m_a}$
- **Then:** $P(|\mu - \hat{\mu}_m| \geq b) \leq 2T^{-4}$ which converges to zero very quickly:
 - Notice:
 - If we don't play action a , its upper bound b increases
 - This means we never permanently rule out an action no matter how poorly it performs
 - Prob. our upper bound is wrong decreases with time T


UCB₁ Algorithm

■ UCB1 (Upper confidence sampling) algorithm

- Set: $\widehat{\mu}_1 = \dots = \widehat{\mu}_k = 0$ and $m_1 = \dots = m_k = 0$

- $\widehat{\mu}_a$ is our estimate of payoff of arm a
 - m_a is the number of pulls of arm a so far

Upper confidence interval (Hoeffding's inequality)



- For $t = 1:T$

- For each arm a calculate: $UCB(a) = \widehat{\mu}_a + \alpha \sqrt{\frac{2 \ln t}{m_a}}$
 - Pick arm $j = \arg \max_a UCB(a)$
 - Pull arm j and observe y_t
 - Set: $m_j \leftarrow m_j + 1$ and $\widehat{\mu}_j \leftarrow \frac{1}{m_j} (y_t + (m_j - 1) \widehat{\mu}_j)$

α ...is a free parameter trading off exploration vs. exploitation

UCB₁: Discussion

■ $UCB(a) = \widehat{\mu}_a + \alpha \sqrt{\frac{2 \ln t}{m_a}}$

$$b \geq \sqrt{\frac{\ln\left(\frac{2}{\delta}\right)}{2m}}$$

- Confidence interval **grows** with the total number of actions t we have taken
- But **shrinks** with the number of times m_a we have tried arm a
- This ensures each arm is tried infinitely often but still balances exploration and exploitation
- α plays the role of δ : $\alpha = f\left(\frac{2}{\delta}\right)$

$$P(|\mu - \widehat{\mu}_m| \geq b) = \delta$$

“Optimism in face of uncertainty”:


The algorithm believes that it can obtain extra rewards by reaching the unexplored parts of the state space

Summary so far

- k -armed bandit problem as a formalization of the exploration-exploitation tradeoff
- Analog of online optimization (e.g., SGD, BALANCE), but with **limited feedback**
- **Simple algorithms are able to achieve no regret (in the limit)**
 - Epsilon-greedy
 - UCB (Upper Confidence Sampling)

Example

- 10 actions, 1M rounds, uniform $[0,1]$ rewards



Use-case: Pinterest

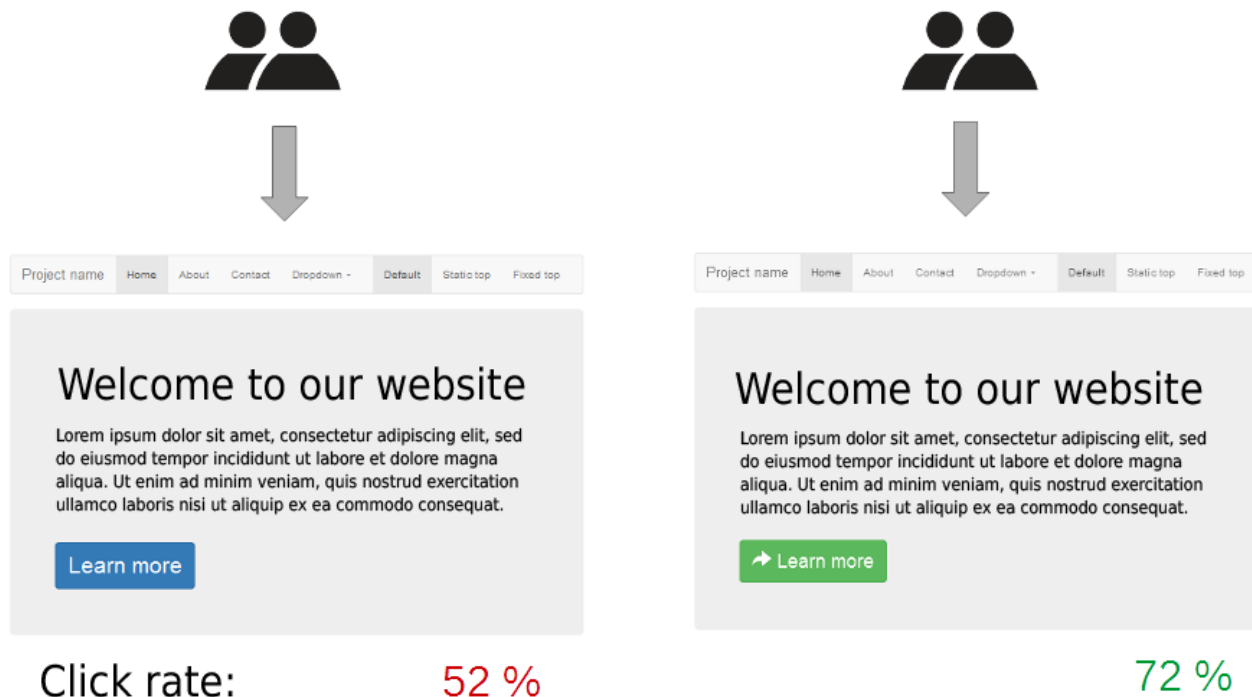
- **Problem:** For new pins/ads we do not have enough signal on how good they are
 - How likely are people to interact with them?
- **Idea:**
 - Try to maximize the rewards from several unknown slot machines by deciding which machines and the order to play
 - Each pin is regarded as an arm, user engagement are considered as rewards
 - Making tradeoff between exploration and exploitation, avoid keep showing the best known pins and trapping the system into local optima

Use-case: Pinterest

- **Solution: Bandit algorithm in round t**
 - **(1)** Algorithm observes user is seeing a set \mathbf{A} of pins/ads
 - **(2)** Based on payoffs from previous trials, algorithm chooses arm $\mathbf{a} \in \mathbf{A}$ and receives payoff $r_{t,\mathbf{a}}$
 - Note only feedback for the chosen \mathbf{a} is observed
 - **(3)** Algorithm improves arm selection strategy with each observation $(\mathbf{a}, r_{t,\mathbf{a}})$
- **If the score for a pin is low, filter it out**

Use Case: A/B testing

- A/B testing is a controlled experiment with two variants, A and B
- Part of the traffic sees variant A, part variant B



Use Case: A/B testing

- Part of the traffic sees variant A, part variant B
- Hypothesis test: does variant A outperform variant B? What test to perform?

Assumed Distribution	Example	Standard Test
Gaussian	Average Revenue Per Paying User	Welch's t-test (Unpaired t-test)
Binomial	Click Through Rate	Fisher's exact test
Poisson	Transactions Per Paying User	E-test
Multinomial	Number of each product purchased	Chi-squared test

- If **A** outperforms **B**, we want to stop the experiment as soon as possible

Use Case: A/B testing

- Imagine you have two versions of the website and you'd like to test which one is better
 - Version **A** has engagement rate of **5%**
 - Version **B** has engagement rate of **4%**
- You want to establish with **95% confidence** that **version A is better**
 - Using t-test, you'd need 22,330 observations (11,165 in each arm) to establish that
- Can bandits do better?

Example: Bandits vs. A/B testing

- **How long does it take to discover $A > B$?**
 - **A/B test:** We need 22,330 observations. Assuming 100 observations/day, we need 223 days
- **The goal is to find the best action (A vs. B)**
- The randomization distribution (traffic to A vs. B) can be updated as the experiment progresses
- **Idea:**
 - Twice per day, examine how each of the variations/arms has performed
 - Adjust the fraction of traffic that each arm will receive going forward
 - An arm that appears to be doing well gets more traffic, and an arm that is clearly underperforming gets less

Thompson Sampling

- **Thompson sampling** assigns sessions to arms in proportion to the probability that each arm is optimal.
- Assume outcome distribution in the set $\{0,1\}$
 - The arm either converts or not
- Then we flip a coin with probability $\theta \rightarrow$ Bernoulli distribution!
- To estimate θ , we count up numbers of ones and zeros


Thompson Sampling: Bernoulli Case

- Given observed 1s and 0s, how do we calculate the distribution of possible values of θ ?
- **Let:**
 - $\theta = (\theta_1, \theta_2, \dots, \theta_k)$... the vector of conversion rates for arms $1, \dots, k$.
 - $\theta_i = \text{\#successes} / (\text{\#successes} + \text{\#failures})$

Beta-Bernoulli Case

- $\text{Beta}(\alpha, \beta) \rightarrow$ Given a 0's and b 1's, what is the distribution over means?

- Prior \rightarrow pseudocounts



- Likelihood \rightarrow observed counts
- Posterior \rightarrow pseudocounts + observed counts

Thompson Sampling

- Arm probabilities θ can be computed using sampling:
 - Each element of θ is an independent random variable from a Beta distribution ($\alpha + \text{successes}, \beta + \text{failures}$)

Algorithm 2 Thompson sampling for the Bernoulli bandit

Require: α, β prior parameters of a Beta distribution

$S_i = 0, F_i = 0, \forall i. \{\text{Success and failure counters}\}$

for $t = 1, \dots, T$ **do**

for $i = 1, \dots, K$ **do**

 Draw θ_i according to $\text{Beta}(S_i + \alpha, F_i + \beta)$.

end for

 Draw arm $\hat{i} = \arg \max_i \theta_i$ and observe reward r

if $r = 1$ **then**


$S_{\hat{i}} = S_{\hat{i}} + 1$

else

$F_{\hat{i}} = F_{\hat{i}} + 1$

end if

end for



Thompson Sampling in General

Thompson Sampling:

- 1. Specify prior (in Beta case often $\text{Beta}(1,1)$)
- 2. Sample from each posterior distribution to get estimated mean for each arm
- 3. Pull arm with highest mean
- 4. Repeat step 2 & 3 forever

Back to Our Problem

But, in our case we have to set the amount of traffic. Set it to be proportional to $P(I_a)$:

- (1) Simulate many draws from $Beta(\alpha + S_a, \beta + F_a)$:

Time	Arm 1	Arm 2	Arm 3
1	0.54	0.73	0.74
2	0.55	0.66	0.73
3	0.53	0.81	0.80
...			

- (2) The probability that arm a is optimal is the empirical fraction of rows for which arm a had the largest simulated value
- (3) Set traffic to arm a to be equal to % of wins of arm a

Reminder: Use Case


- **Imagine you have two versions of the website and you'd like to test which one is better**
 - Version **A** has engagement rate of **5%**
 - Version **B** has engagement rate of **4%**
- **You want to establish with 95% confidence that version A is better**
 - You'd need 22,330 observations (11,165 in each arm) to establish that
 - Use t-test to establish the sample size
- **Can bandits do better?**

Example

A/B test: We need 22,330 observations. Assuming 100 observations/day, we need 223 days

- On 1st day about 50 sessions are assigned to each arm
- Suppose **A** got really lucky on the first day, and it appears to have a 70% chance of being superior
- Then we assign it 70% of the traffic on the second day, and the variant B gets 30%
- At the end of the 2nd day we accumulate all the traffic we've seen so far (over both days), and recompute the probability that each arm is best

Simulation



- The experiment finished in 66 days, so it saved you 157 days of testing (66 vs 223)

Generalization to multiple arms

- Easy to generalize to multiple arms:

